



# MONGODB VS COUCHBASE SERVER

Sreeja Matturu  
Jithin Jacob Benjamin Jacob

# MongoDB and Couchbase

- ▶ NoSQL Distributed database
- ▶ Document oriented
- ▶ Share-Nothing architecture
- ▶ Stores JSON type data

# Introduction

- ▶ Taken two huge JSON datasets called as “Business” and “Tip” from “Yelp” dataset
- ▶ Yelp dataset contains the information about different businesses in different cities and states.
- ▶ Business dataset has 192K documents that in it, while the Tip dataset had 1.2M documents in it.
- ▶ Imported the datasets into both the systems and tested on 1 master with 3 slave, 1 master with 1 slave and a single 1 master setup.

# Installation of MongoDB

- ▶ Implemented both the systems using Amazon Web Services (AWS)
- ▶ Installed MongoDB Atlas, which is the cloud version of MongoDB with m5.a large EC2 instances
- ▶ Primarily created four clusters determining them as one master and 3 slaves
- ▶ Connected the MongoDB Atlas to the terminal and performed different queries within the terminal.
- ▶ The performance was evaluated from the MongoDB Atlas interface.

# Couchbase Installation

- ▶ Couchbase Server Enterprise Edition and Couchbase Sync Gateway was installed on AWS EC2 instances
- ▶ For Good Replication and High Availability, 1 Couchbase sync gateway instances is required for 2 Couchbase servers
- ▶ Created a new cluster With 6 instances where 1 is the master and 3 are slaves and 2 systems are Couchbase sync gateway servers
- ▶ Primary design for the Couchbase is the one master with 3 slaves
- ▶ Dataset import was made by using cbimport through the terminal of the local machine

# Hardware Specification

- ▶ 4 instances for MongoDB and 6 instances for Couchbase server taken from AWS
- ▶ Each instance has:
  - ▶ 8GB of RAM
  - ▶ 2-core vCPU – Intel Xeon Platinum 8000 series processor
  - ▶ Up to 10Gbps of Network Bandwidth
  - ▶ Elastic Block Storage bandwidth of 2120Mbps
  - ▶ All chosen from AWS us-east-1 (North Virginia)

# Query Scenario


- ▶ Point Query to show the result of the name of businesses from the Business dataset that has the review\_count>1000
- ▶ Count aggregation with sort to count each compliment\_count value in the Tip dataset and sort them in ascending order.
- ▶ Join Business and Tip dataset using a unique ID called business\_id.
- ▶ Min and Max aggregation operation to show the highest and lowest review\_count from Business dataset.
- ▶ To calculate the review\_count between 0 and 50 in the Business dataset.
- ▶ Map Reduce to calculate the review\_count range between 0 and 50.

# Query Performance Analysis


- ▶ The performance analysis for both the systems was made on:
  - ▶ One Master + Three Slaves
  - ▶ One Master + One Slave
  - ▶ Only One Master



## 1 Master - 3 Slave





## 1 Master - 1 Slave



# 1 Master

Plot Area



- 
- ▶ How are multiple queries executed parallelly?
    - ▶ Allows multiple users to read and write the same data.
    - ▶ Uses Lock
  - ▶ How are single queries parallelized?
    - ▶ Intra-operation Parallelism is followed by MongoDB
  - ▶ What is the sort algorithm?
    - ▶ Uses In-Memory Sort algorithm
    - ▶ Uses Top k-sort algorithm
  - ▶ How are Joins parallelized?
    - ▶ Intra Query parallelism
  - ▶ Are query plans evaluated parallelly?
    - ▶ Multiple indexes are evaluated in parallel

# Conclusion

- ▶ The cost for the Couchbase is higher when compared to MongoDB in AWS.
- ▶ From the query execution analysis, MongoDB is faster in point queries and aggregate functions but not the join operations.
- ▶ The performance for range queries can be increased by creating more indexes manually in MongoDB.
- ▶ In MongoDB, one master with three slave has higher performance than one with single master.





Thank  
you!!